

# Noise-Robust Bird Vocalization Classification for Low-Resource Scenarios: A Multi-Paradigm Comparative Study

1<sup>st</sup> Jianan Zhang  
Department of Computer Science  
Stevens Institute of Technology  
Beijing, China  
email: jzhang112@stevens.edu

2<sup>nd</sup> Jincheng Chen  
Department of Computer Science  
Stevens Institute of Technology  
Guangzhou, China  
email: jchen174@stevens.edu

3<sup>rd</sup> Wenjuan Huang  
Department of Computer Science  
Stevens Institute of Technology  
Zhejiang, China  
email: whuang32@stevens.edu

**Abstract**—Automated bird vocalization classification from passive acoustic monitoring is critical for biodiversity conservation, yet field recordings present two compounding challenges: environmental noise that masks subtle species-specific features, and extreme long-tail distributions where rare species have minimal training samples. Existing solutions such as BirdNET and Google Perch lack standardized noise robustness evaluation and do not provide cross-paradigm comparisons incorporating traditional machine learning baselines. This work presents a systematic multi-paradigm comparative study on a unified dataset of 129,834 recordings spanning 1,126 bird species from BirdCLEF 2021–2026. We evaluate three approaches: (1) frozen YAMNet with a Dense head on 3072-dimensional clip-level pooled embeddings (mean, max, and std pooling over 1024-dimensional frame-level YAMNet outputs every 0.48 s; encoder not fine-tuned); (2) LightGBM on the same 3072-dimensional pooled features—a cascade decoupling neural feature extraction from tree-based classification, achieving faster inference (7.7 ms/s) than BirdNET (~10.0 ms/s) with modular flexibility; and (3) FastAI ConvNeXt-Small on mel-spectrograms. All share identical 5-fold CV splits, the same noise-evaluation protocol, CSV `loss_class_weight`, and targeted audio augmentation for low-resource species (<15 samples). Noise robustness is evaluated under Gaussian white noise at 5 dB, 0 dB, and −5 dB SNR. Our primary contribution is a reproducible evaluation framework with common data splits and noise testing for direct quantitative comparison across heterogeneous ML paradigms—a capability absent from existing bioacoustic tools.

## I. INTRODUCTION

Automated bird vocalization classification is essential for large-scale biodiversity monitoring, yet field recordings pose dual challenges: environmental noise masks discriminative spectral features, and species distributions follow an extreme long-tail pattern where rare species have minimal training data [1]–[3]. These compounding factors severely degrade classifier performance, particularly for the very species most in need of conservation monitoring.

Existing off-the-shelf solutions such as BirdNET [1] and Google Perch [2] demonstrate significant limitations under these conditions. BirdNET’s accuracy degrades markedly when SNR falls below 3 dB [3], and its recall varies drastically across regions—from PR AUC 0.16–0.23 in North America and Europe to 0.03–0.04 in Africa and Asia [3]. Perch achieves only ~60% accuracy on BirdCLEF+ 2025 data containing

species outside its training set [2]. Critically, neither provides standardized, quantifiable noise robustness evaluation nor cross-paradigm comparisons that include traditional ML baselines.

This project addresses these gaps through three pipelines on a unified 129,834-sample dataset with controlled noise evaluation: (1) frozen YAMNet (MobileNetV1, AudioSet [4]) with a Dense head on 3072-dimensional clip-level pooled embeddings (aggregated from 1024-dimensional frame-level outputs); (2) LightGBM on the same 3072-dimensional pooled features—a two-stage cascade decoupling neural feature extraction from tree-based classification, enabling GPU-free training and modular experimentation; and (3) FastAI ConvNeXt-Small on mel-spectrograms. All share identical 5-fold CV, the same SNR evaluation protocol, CSV `loss_class_weight`, and targeted audio augmentation for species with fewer than 15 training samples. Evaluation uses clean audio plus Gaussian noise at 5 dB, 0 dB, and −5 dB SNR. While we use Gaussian white noise as a controlled, reproducible baseline, qualitatively different field noise (wind, rain, insect interference) is discussed in Section V.

Our contributions are threefold: (i) a reproducible cross-paradigm evaluation framework with common data splits and standardized noise testing; (ii) empirical insights into trade-offs between frozen embedding classifiers, embedding-based gradient boosting, and spectrogram-based classification—with emphasis on the YAMNet+LightGBM cascade as a computationally efficient alternative to monolithic systems; and (iii) a practical framework combining shared CSV-based class weighting and targeted low-resource augmentation with standardized evaluation for low-resource bioacoustic classification.

## II. RELATED WORK

### A. BirdNET and Google Perch: Off-the-Shelf Solutions

BirdNET [1] (Cornell Lab/Chemnitz University) uses an EfficientNetB0-like backbone with 1024-dimensional embeddings, classifying 48 kHz audio into 3,337 classes [5]. Despite its widespread adoption, it exhibits: (1) severe performance degradation below 3 dB SNR [3]; (2) highly heterogeneous recall across regions (0.24–0.52) and species rarity levels [3]; and (3) label noise and dataset imbalances as primary error

sources [6]. Google Perch [2] (EfficientNet-B1) achieves state-of-the-art on BirdSET and BEANS benchmarks [7], [8] but only  $\sim 60\%$  accuracy on BirdCLEF+ 2025 data with out-of-training-set species. Critically, neither provides standardized controllable noise testing pipelines.

### B. BirdCLEF Competitions: Persistent Challenges

The BirdCLEF series (2021–2026) consistently benchmarks core challenges: long-tailed distributions, weak labeling (recording-level vs. segment-level), and domain shift between Xeno-Canto training and PAM test soundscapes [9], [10]. BirdCLEF+ 2025 saw foundation models like BEATs achieve 0.913 F1 with pseudo-labeling [11]. These competitions underscore that no integrated solution spans data, model, and evaluation dimensions for low-resource noisy environments.

### C. Noise Robustness and Low-Resource Classification

Environmental noise remains a formidable obstacle. Noise reduction preprocessing can help but occasionally degrades performance by removing informative features [12]. Data augmentation with real environmental noise under varying SNR has proven effective [9], [13], and generative approaches (ACGANs, DDPMs) show promise [14]. For low-resource species, transfer learning is essential: avian-specific embeddings outperform general audio embeddings [15], and AudioSet pre-training followed by fine-tuning boosts low-resource performance [15], [16]. Few-shot approaches (Siamese networks) target rare species absent from BirdNET/Perch [17]. However, existing studies rarely provide systematic cross-paradigm comparisons on identical datasets.

### D. YAMNet and Audio Transfer Learning

YAMNet [4], [16] (MobileNetV1, AudioSet-trained) outputs 1024-dimensional frame-level embeddings at 0.48-second intervals. In our pipelines, these frames are aggregated by mean, maximum, and standard deviation into 3072-dimensional clip-level pooled features for downstream classifiers. It has been adapted for speech recognition, urban noise, and industrial monitoring [18]. While domain-specific models (Perch, BirdNET) outperform general-purpose ones on complex calls [15], [19], YAMNet provides a computationally accessible baseline for raw audio classification.

### E. Research Gap Summary

The literature reveals six critical gaps: (1) no unified framework comparing tabular ML, image CNN, and raw-audio transfer learning on identical datasets with standardized noise evaluation; (2) off-the-shelf tools lack quantifiable controllable noise pipelines; (3) no systematic cross-paradigm comparison for rare species; (4) no reproducible controlled-variable experiment separating clean baseline from gradient noise; (5) BirdCLEF solutions rarely integrate data, model, and evaluation dimensions; and (6) no systematic evaluation of hybrid pipelines combining deep embeddings with lightweight tree classifiers—which can outperform monolithic systems while being faster and more deployable. Our YAMNet+LightGBM cascade directly addresses gap (6), achieving  $\sim 7.7$  ms/s embedding latency (faster than BirdNET’s  $\sim 10.0$  ms/s) with modular flexibility.

TABLE I: Final Dataset Split (V2.0)

| Split                    | Records       | Species | Use                                           |
|--------------------------|---------------|---------|-----------------------------------------------|
| Cleaned population       | 144,250       | 1,126   | Source                                        |
| Training pool            | 129,834       | 1,126   | Development & CV                              |
| CV validation (per fold) | $\sim 25,967$ | 1,126   | In-fold validation (5-fold SNR-stratified CV) |
| Fixed holdout test       | 14,416        | 1,093   | Final evaluation                              |

TABLE II: SNR Tier Distribution Across Train/Test Splits

| SNR Tier            | Training | Test  |
|---------------------|----------|-------|
| High ( $\geq 4.0$ ) | 86,731   | 9,631 |
| Medium (2.0–3.9)    | 41,011   | 4,553 |
| Heavy ( $< 2.0$ )   | 2,092    | 232   |

## III. OUR SOLUTION

### A. Dataset

1) *Data Source and Collection*: Compiled from BirdCLEF 2021–2026, aggregating Xeno-Canto and iNaturalist recordings [9], [20]. Raw merge of six annual CSVs: 183,239 records. Field-level deduplication (zero-information-loss merge) yielded 167,308 unique recordings.

2) *Data Cleaning and Quality Control*: Two filters applied:

- 1) Remove numeric-only `primary_label` (non-avian taxa: Amphibia, Insecta, Mammalia, Reptilia)—excludes 101 non-bird species;
- 2) Discard rating  $\leq 0.5$  (unrated/extremely low-quality).

Final cleaned population: **144,250 recordings across 1,126 bird species.**

SNR proxy tier derived from rating: High ( $\geq 4.0$ ), Medium (2.0–3.9), Heavy ( $< 2.0$ ). Split: 90/10 with SNR stratification (seed=42). Species with  $< 5$  recordings retained entirely in training. Tables I–II summarize the split.

5-fold SNR-stratified CV: each fold  $\sim 103,853$  training,  $\sim 25,967$  validation. No filename overlap between train/test or train/val within folds. This represents a  **$10\times$  scale-up** from the original proposal (12,400  $\rightarrow$  129,834), enabling robust long-tail evaluation.

3) *Long-Tail Distribution and Mitigation*: Species distribution: median  $\sim 72$  training samples/species, minimum 1, head class 1,300. **Shared long-tail mitigation (all three pipelines)**:

- 1) **Balanced sampling**:  $w_{\text{sampl}} = 1/N_c$ , ensuring equal per-species batch probability (exported for LightGBM; not used in Dense-head training).
- 2) **Mild class-weighted loss**:

$$w_{\text{loss}}(c) = \text{clip}\left(\frac{1/\sqrt{N_c}}{\mu_w}, 0.25, 4.00\right) \quad (1)$$

Inverse-square-root is milder than  $1/N_c$ ; clipping ensures stability. Species with fewer than 15 training samples are treated as low-resource tail classes, consistent with BirdCLEF findings that such species exhibit near-random performance [10].

- 3) **Targeted audio augmentation:** 159 species (14.1%) with <15 training examples receive four on-the-fly waveform transformations implemented in `audio_augmentation.py`: time stretching ( $0.85\times-1.10\times$ ), pitch shifting ( $\pm 1\pm 2$  semitones), additive Gaussian noise (5–15 dB SNR), and volume scaling ( $0.70\times-1.30\times$ ), capped at 16 variants per recording and 50 new variants per species, with a per-species target of 15 samples. This augmentation is enabled in all three reported pipelines: YAMNet applies `expand_with_augmentation` before embedding extraction and Dense-head training; LightGBM trains on 3072-dimensional pooled embeddings computed from the same augmented training folds; FastAI generates augmented mel-spectrogram PNGs via `augmentation_glue.py` before ConvNeXt training.

4) *Evaluation Protocol:* Fixed test set (14,416) reserved for final evaluation only. Metrics:

- 1) clean Top-1 accuracy;
- 2) noise robustness at 5 dB, 0 dB, -5 dB SNR (Gaussian white noise, deterministic per-sample seeding with base seed=42, peak-normalized);
- 3) per-SNR-tier stratified accuracy;
- 4) inference latency (ms/sample) and peak GPU memory.

## B. Machine Learning Algorithms

1) *Frozen YAMNet Embeddings with Dense Head:* **Rationale.** MobileNetV1 pre-trained on AudioSet’s 521 categories [4] provides general-purpose acoustic features transferable to bird vocalizations. AudioSet includes environmental noise [16], giving inherent noise-robust representations. We freeze the encoder and train only a lightweight head on pooled embeddings [15].

**Design.** 16 kHz mono, 5-second clips (center-crop or zero-pad), peak-normalized. (1) Frozen YAMNet encoder (`trainable=False`) outputs 1024-dimensional frame-level embeddings every 0.48 s; mean, max, and std pooling across frames  $\rightarrow$  3072-dimensional clip-level feature vector. (2) Dense head: Dense(256, ReLU)  $\rightarrow$  Dropout(0.3)  $\rightarrow$  Dense(1126, Softmax). Adam ( $\text{lr} = 10^{-3}$ ), up to 50 epochs, early stopping (`patience=8`), ReduceLROnPlateau (`factor=0.5`, `patience=4`), `loss_class_weight` as training sample weights. Low-resource species (<15 samples) are expanded via `expand_with_augmentation` before embedding extraction. Embeddings are cached once and exported per fold for LightGBM. Noise evaluation re-computes clip-level pooled embeddings on noisy test audio.

2) *LightGBM: Gradient-Boosted Decision Trees on YAMNet Embeddings:* **Rationale.** Gradient-boosted trees complement deep learning: no GPU required, fast CPU training, native feature importance, robust to hyperparameters. However, trees cannot operate on raw audio. Our **cascade architecture** bridges this: YAMNet’s frozen encoder extracts embeddings; LightGBM classifies. This decoupling offers three advantages: (i) embedding extraction runs once per file; (ii) classification is GPU-free, trainable in minutes with class-weighted loss; (iii) modular—same embeddings reusable for different classifiers.

**Audio Preprocessing.** 16 kHz mono OGG loaded directly via `librosa` (no WAV conversion). Fixed 5.0 seconds (zero-pad if shorter; center-crop if longer), peak-normalized to  $[-1, 1]$ . **No noise reduction**—evaluation requires measuring robustness against controlled additive noise.

**Feature Extraction.** The frozen YAMNet encoder produces 1024-dimensional frame-level embeddings at 0.48 s intervals ( $\sim 10$  frames per 5 s clip). These are pooled by **mean, maximum, and standard deviation** into a **3072-dimensional clip-level feature vector**—not the raw 1024-d frame outputs. This captures both average spectral profile and temporal variability—critical for dynamic vocalization patterns. Compared to handcrafted features (MFCC statistics, spectral centroid), YAMNet pooled features preserve deeper semantic information; compared to flattened spectrograms ( $\sim 64,000$  dims), 3072 dims makes tree learning tractable.

**Model Design.** LightGBM: 1,500 rounds with early stopping (`patience=50`), learning rate 0.05, `max_depth=6`, `num_leaves=31`, L1/L2 regularization ( $\alpha = 0.1$ ,  $\lambda = 0.1$ ), `loss_class_weight` as sample weights. CPU inference only. Outputs 1,126-dimensional raw class scores under the multiclass objective (margin logits, not native softmax). For probability-based metrics, scores are normalized via softmax at evaluation time.

**Methodological value:** Establishes performance floor for CPU-only tools with high-quality embeddings—quantifying accuracy-per-compute trade-offs against deep neural classifiers on the same embeddings.

3) *FastAI ConvNeXt: Mel-Spectrogram Image Classification:* **Rationale.** Mel-spectrograms expose harmonic structures, formant patterns, and temporal modulations. ConvNeXt-Small provides strong performance with ImageNet transfer learning [10], [13]. Long-tail imbalance is handled with the same CSV `loss_class_weight` as YAMNet/LightGBM, supplemented by spectrogram-level MixUp and RandomErasing. Low-resource waveform augmentation is applied upstream via `augmentation_glue.py` before PNG caching.

**Design.** 16 kHz mono; highest-RMS 5 s segment; 128 mel bands (`n_fft=1024`, hop 512, 500–8000 Hz)  $\rightarrow$  Viridis RGB images at  $224\times 224$  (cached offline, including augmented low-resource samples). ConvNeXt-Small, FP16, CSV `loss_class_weight` as sample weights, MixUp ( $\alpha = 0.4$ ), RandomErasing ( $p = 0.4$ ). `fine_tune`: 3 frozen epochs + 12 full-network epochs (`base_lr=2\times 10^{-3}`). Metrics: Top-1 and macro-F1.

## C. Implementation Details

All pipelines share:

- Identical 5-fold CV splits (SNR-stratified)
- Identical noise evaluation: Gaussian white noise at 5/0/-5 dB, deterministic per-sample seeding (base seed=42)
- Same metrics: Top-1, Top-5, macro-F1, balanced accuracy, stratified by SNR tier and class frequency group
- Inference measurement: 5 warm-up + 50 timed iterations; GPU memory profiling for FastAI fine-tuning

(YAMNet embedding extraction uses <1 GB as a one-time cost); CPU-only for LightGBM

- Shared long-tail controls: all pipelines use CSV `loss_class_weight` and targeted audio augmentation for species with <15 samples; FastAI additionally applies MixUp and RandomErasing at the spectrogram level
- Embedding terminology: YAMNet outputs 1024-d frame-level embeddings; Dense/LightGBM consume 3072-d clip-level pooled features (mean/max/std)

#### IV. COMPARISON

##### A. Experimental Results

[To be completed after model training and evaluation. Will present: (1) comparative Top-1, Top-5, macro-F1 on clean test audio; (2) noise robustness degradation curves (5 dB, 0 dB, -5 dB); (3) per-SNR-tier and per-class-frequency-group stratified performance; (4) inference latency and GPU memory benchmarks.]

##### B. Comparison with Existing Solutions

[To be completed. Will compare best-performing model against BirdNET and Perch baselines on the same test set, analyzing trade-offs between transfer learning, YAMNet embedding-based gradient boosting, and CNN image classification.]

#### V. FUTURE DIRECTIONS

- 1) **Multi-label classification:** Incorporate `secondary_labels` for partially labeled co-occurrences, reducing label noise and improving tail-class performance.
- 2) **Attention-based embedding aggregation:** Replace simple mean/max/std pooling with learnable attention to preserve temporally localized discriminative patterns.
- 3) **Extended noise conditions:** Evaluate against wind, rain, and insect noise—qualitatively different from Gaussian.
- 4) **Model distillation:** Distill YAMNet+LightGBM cascade into a single lightweight edge-deployable model.
- 5) **Active learning for rare species:** Identify most informative unlabeled recordings for human annotation, maximizing labeling efficiency for species with <5 samples.

#### VI. CONCLUSION

[To be completed after experimental results are obtained. This section will summarize: whether the noise-robust bird vocalization classification problem is adequately addressed by our multi-paradigm approach; which algorithm(s) demonstrate superior performance under specific conditions (clean vs. noisy, head vs. tail classes); which shared techniques (CSV `loss_class_weight`, targeted low-resource augmentation, frozen AudioSet embeddings vs. ConvNeXt spectrogram modeling) contribute most to performance gains; and broader implications for low-resource bioacoustic classification system design.]

#### ACKNOWLEDGMENT

The authors would like to thank their advisor for guidance and the Kaggle community for providing the BirdCLEF datasets.

#### REFERENCES

- [1] S. Kahl, C. M. Wood, M. Eibl, and H. Klinck, “BirdNET: A deep learning solution for avian diversity monitoring,” *Ecological Informatics*, vol. 61, p. 101236, 2021.
- [2] B. van Merriënboer *et al.*, “Perch 2.0: The bittern lesson for bioacoustics,” *arXiv preprint arXiv:2508.04665*, 2025. [Online]. Available: <https://arxiv.org/abs/2508.04665>
- [3] D. Funosas *et al.*, “A global assessment of BirdNET performance: Differences among continents, biomes, and species,” *Ecological Indicators*, vol. 182, p. 114550, 2026.
- [4] S. Hershey, S. Chaudhuri, D. P. W. Ellis, J. F. Gemmeke, A. Jansen, R. C. Moore, M. Plakal, D. Platt, R. A. Saurous, B. Seybold, M. Slaney, R. Weiss, and K. Wilson, “CNN architectures for large-scale audio classification,” in *Proc. IEEE Int. Conf. Acoust., Speech Signal Process. (ICASSP)*, 2017.
- [5] S. Kahl, C. M. Wood, and H. Klinck, “BirdNET model v2.4 (EfficientNetB0-like backbone, 1024-dimensional embedding),” Zenodo, 2025.
- [6] M. Assarzadeh, P. Atufi, M. Eibl, and S. Kahl, “Investigating misclassification in bird sounds: The adverse effect of label noise on BirdNET,” *Ecological Informatics*, vol. 92, p. 103494, 2025.
- [7] L. Rauch *et al.*, “BirdSet: A large-scale dataset for audio classification in avian bioacoustics,” *arXiv preprint arXiv:2403.10380*, 2024. [Online]. Available: <https://arxiv.org/abs/2403.10380>
- [8] M. Hagiwara *et al.*, “BEANS: The benchmark of animal sounds,” in *Proc. IEEE Int. Conf. Acoust., Speech Signal Process. (ICASSP)*, 2023, arXiv:2210.12300. [Online]. Available: <https://arxiv.org/abs/2210.12300>
- [9] Cornell Lab of Ornithology, “BirdCLEF / BirdCLEF+ competitions (2021–2026),” Kaggle, 2026. [Online]. Available: <https://www.kaggle.com/competitions/birdclef-2026>
- [10] E. Witting, H. de Heer, J. Lim, C. T. Kopar, and K. Sándor, “Addressing the challenges of domain shift in bird call classification for BirdCLEF 2024,” in *CEUR Workshop Proc.*, vol. 3740, 2024.
- [11] S. Chen *et al.*, “BEATS: Audio pre-training with acoustic tokenizers,” in *Proc. Int. Conf. Mach. Learn. (ICML)*, 2023, arXiv:2212.09058. [Online]. Available: <https://arxiv.org/abs/2212.09058>
- [12] T. Denton, S. Wisdom, and J. R. Hershey, “Improving bird classification with unsupervised sound separation,” in *Proc. IEEE Int. Conf. Acoust., Speech Signal Process. (ICASSP)*, 2022, arXiv:2110.03209. [Online]. Available: <https://arxiv.org/abs/2110.03209>
- [13] A. Miyaguchi, N. Zhong, M. Gustineli, and C. Hayduk, “Transfer learning with semi-supervised dataset annotation for birdcall classification,” in *CLEF 2023 Working Notes*, 2023.
- [14] A. Gibbons, E. King, I. Donohue, and A. Parnell, “Generative AI-based data augmentation for improved bioacoustic classification in noisy environments,” *arXiv preprint arXiv:2412.01530*, 2024. [Online]. Available: <https://arxiv.org/abs/2412.01530>
- [15] B. Ghani, T. Denton, S. Kahl, and H. Klinck, “Global birdsong embeddings enable superior transfer learning for bioacoustic classification,” *Scientific Reports*, vol. 13, p. 22876, 2023.
- [16] J. F. Gemmeke *et al.*, “Audio set: An ontology and human-labeled dataset for audio events,” in *Proc. IEEE Int. Conf. Acoust., Speech Signal Process. (ICASSP)*, 2017. [Online]. Available: <https://research.google/pubs/audio-set-an-ontology-and-human-labeled-dataset-for-audio-events/>
- [17] M. Zhong, J. LeBien, M. Campos-Cerqueira, T. M. Aide, R. Dodhia, and J. Lavista Ferres, “A comparison of convolutional neural networks and few-shot learning in classifying long-tailed distributed tropical bird songs,” *bioRxiv*, 2023.
- [18] L. Liu, Q. Xu, S. Mao, J. Mu, X. Zhao, W. Song, and P. Cheng, “YAMNet-based transfer learning for compact noise classification in urban and wireless systems,” *EURASIP Journal on Wireless Communications and Networking*, vol. 2025, no. 1, p. 74, 2025.

- [19] Google Research, “VGGish: Audio feature embedding model,” Kaggle Models, 2024. [Online]. Available: <https://www.kaggle.com/models/google/vggish>
- [20] Xeno-Canto Foundation, “Xeno-Canto: Sharing bird sounds from around the world,” <https://xeno-canto.org/>, 2024.